

Utilizing the R Statistical Language for Analyzing Server Logs

Ryan Talabis and Robert McPherson

January 26, 2013

Outline

Processing the Data

Date and Time Stamp Handling

Text Analysis

Outlier Detection

Processing the Data

Read in the Data

- Data can be read from text files of server logs.
- An R function called scan is useful for this.
- Corpus function is used to convert data into a corpus, consisting of documents and words.
- Each line in the log can be considered a document, so that comparisons can be made across time.

```
> #Read server log - text file
> secureData <- scan("secure.txt", character(0), sep = "\n") # separate each line
> #Save as a word corpus
> data <- Corpus(VectorSource(secureData[1:1000],))
> inspect(data[1:5]) #Inspect the first five entries in the data
```

A corpus with 5 text documents

The metadata consists of 2 tag-value pairs and a data frame Available tags are:

create_datercreatorAvailablevariablesinthedataframeare : MetaID

[[1]] Nov 10 07:53:34 uiftp sshd[28959]:

pam_unix(sshd : session) : sessionclosedforuserrichard

[[2]] Nov 10 08:30:26 uiftp sshd[17793]: Accepted publickey for richard from 10.130.192.162 port 34178 ssh2

[[3]] Nov 10 08:30:26 uiftp sshd[17793]:

pam_unix(sshd : session) : sessionopenedforuserrichardby(uid = 0)

[[4]] Nov 10 08:30:26 uiftp sshd[17793]:

pam_unix(sshd : session) : sessionclosedforuserrichard

[[5]] Nov 10 08:30:29 uiftp sshd[17819]: Accepted publickey for richard from 10.130.192.162 port 34180 ssh2

Function to Trim White Space

- Create a function called, “trim”, as a convenient way to remove white space in the data transformation process.
- Utilizes the R function, gsub, and regular expressions.

```
> #Function to strip white space from strings  
> trim <- function (x) gsub("^\\s+|\\s+$", "", x)
```

Using Regular Expressions to Parse the Time Stamp

- Regular expressions provide a common syntax for parsing characters and text.
- Regular expressions are strings which locate specific patterns in text.
- These regular expressions extract text patterns which identify time, day, and month from the date stamp in the log file.

```
> #Regular expressions for parsing date information
> regexpTime = "[0-9][0-9]:[0-9][0-9]:[0-9][0-9]"
> regexpDay = "\\s[0-9][0-9]\\s"
> regexpMonth = "[A-Z][a-z][a-z]\\s"
> time = trim(str_extract(secureData[1:1000],regexpTime))
> day = trim(str_extract(secureData[1:1000],regexpDay))
> month = trim(str_extract(secureData[1:1000],regexpMonth))
```

Results from Parsing the Time Stamp

- Following is a sample from the results extracted from the time stamp.
- The cbind function means, column bind, and binds arrays into columns.
- The data.frame function puts the columnar data into a single table.

```
> #Make a data frame consisting of month, day, and time columns  
> timeStampCols = data.frame(cbind(month, day, time))  
> timeStampCols[1:5,]
```

	month	day	time
1	Nov	10	07:53:34
2	Nov	10	08:30:26
3	Nov	10	08:30:26
4	Nov	10	08:30:26
5	Nov	10	08:30:29

Transforming the Text Corpus with the tm_map Function

- The tm_map function is from the tm text mining package.
- It maps a number of text mining transformations to the data.
- Transformations include stripping whitespace, changing to lower case, removing common stopwords, removing word suffixes - called stemming, removing punctuation, and removing numbers, among others.

Examples of the tm_map Function

```
> data2 = tm_map(data, stripWhitespace)
> data2 = tm_map(data2, tolower)
> stopWords = c(stopwords("english"), "uiftp", "sshd")
> data2 = tm_map(data2, removeWords, stopWords)
> data2 = tm_map(data2, stemDocument)
> data2 = tm_map(data2, removePunctuation)
> data2 = tm_map(data2, removeNumbers)
> inspect(data2[1:5])
```

A corpus with 5 text documents

The metadata consists of 2 tag-value pairs and a data frame

Available tags are:

create_date creator

Available variables in the data frame are:

MetaID

```
[[1]]
nov      pamunixsession session close  user richard

[[2]]
nov      accept publickey  richard  port  ssh

[[3]]
nov      pamunixsession session  user richard  uid

[[4]]
nov      pamunixsession session close  user richard

[[5]]
nov      accept publickey  richard  port  ssh
```

Making a Document Term Matrix with the DocumentTermMatrix Function

- A document term matrix shows the frequency count of each word.
- Frequencies are shown for each document.
- Documents are in rows.
- Terms, or words, are in columns.

```
> #Make a word frequency matrix, with documents as rows, and terms as columns  
> dtm = DocumentTermMatrix(data2)  
> inspect(dtm[1:5,1:5])
```

A document-term matrix (5 documents, 5 terms)

Non-/sparse entries: 2/23

Sparsity : 92%

Maximal term length: 7

Weighting : term frequency (tf)

	Terms				
Docs	accept	access	address	admin	age
1	0	0	0	0	0
2	1	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0
5	1	0	0	0	0

Removing Sparse Terms

- It can be useful to remove infrequently occurring words.
- Sometimes the most significant terms are the most frequently occurring ones.
- Use the `removeSparseTerms` function.

```
> #Remove and inspect sparse terms, with at least 80% sparse occurrence
> dtm = removeSparseTerms(dtm, 0.85)
> inspect(dtm[1:5,1:5])
```

A document-term matrix (5 documents, 5 terms)

Non-/sparse entries: 10/15

Sparsity : 60%

Maximal term length: 14

Weighting : term frequency (tf)

	Terms				
Docs	close	invalid	nov	pamunixsession	password
1	1	0	1	1	0
2	0	0	1	0	0
3	0	0	1	1	0
4	1	0	1	1	0
5	0	0	1	0	0

Transposing the Term Matrix

- Can opt to have the terms in rows, and the documents in columns.
- Use the TermDocumentMatrix for this.
- Same as the DocumentTermMatrix, but transposed.

```
> #Make a word frequency matrix, with terms as rows, and documents as columns
> dtm2 = TermDocumentMatrix(data2)
> inspect(dtm2[1:5,1:5])
```

A term-document matrix (5 terms, 5 documents)

Non-/sparse entries: 2/23

Sparsity : 92%

Maximal term length: 7

Weighting : term frequency (tf)

	Docs				
Terms	1	2	3	4	5
accept	0	1	0	0	1
access	0	0	0	0	0
address	0	0	0	0	0
admin	0	0	0	0	0
age	0	0	0	0	0

Date and Time Stamp Handling

Append Time Element Columns to Text Frequency Columns

- Convert the word corpus to text, so that it can be appended.
- The example below includes a version with all time stamp elements, and another with only the month.

```
> #Add month, day, and time columns to the document term matrix
> #First, turn the dtm into a plain text object
> dtmText = inspect(dtm)
> #This data frame version has the full month, day, and time stamp
> dtmStamp = data.frame(cbind(timeStampCols, dtmText))
> #This data frame version only has the month appended, not day or time
> dtmMonth = data.frame(cbind(month, dtmText))
```

Utilize SQL as One Means of Grouping Data

- Utilize the sqldf function.
- This function allows for writing SQL queries against an R data frame.

```
> #Sum up the frequencies each of the top terms for each month
> topTermsByMonth = sqldf("select month
, count(*) as rowCount
, sum(user) as user
, sum(invalid) as invalid
, sum(richard) as richard
, sum(port) as port
, sum(password) as password
, sum(ssh) as ssh
, sum(uid) as uid
, sum(session) as session
from dtmMonth group by month")
> topTermsByMonth
```

	month	rowCount	user	invalid	richard	port	password	ssh	uid	session
1	Apr	36	17	4	21	7	2	5	6	17
2	Dec	16	8	0	10	4	0	2	4	8
3	Feb	54	9	0	37	8	33	2	4	8
4	Jun	91	62	26	20	18	14	18	19	15
5	Mar	22	9	0	18	4	5	4	6	9
6	May	611	328	147	103	121	96	117	119	83
7	Nov	170	97	0	46	43	21	43	44	96

Text Analysis

Finding High Frequency Terms

- Frequently occurring words can be identified with the `findFreqTerms` function
- Can select words that occur within a low and high frequency range
- `findFreqTerms(x, lowfreq = 0, highfreq = Inf)`
- Can also specify only a single, lower bound, with upper bound defaulting to infinity

```
> #Find the top 20 most frequent terms  
> findFreqTerms(dtm, 20)
```

[1] "close"	"invalid"	"nov"	"pamunixsession"
[5] "password"	"port"	"richard"	"session"
[9] "ssh"	"uid"	"user"	

Find Associated Terms

- Can find terms in the text that are associated with a given term.
- Use the findAssocs function
- Can specify the minimum correlation.

```
> #Find associated terms with the term, "root", and a correlation of at least 0.4  
> findAssocs(dtm, "session", 0.4)
```

session	pamunixsession	user	close
1.00	0.88	0.49	0.40

```
> findAssocs(dtm, "user", 0.4)
```

user	session	pamunixsession	invalid
1.00	0.49	0.45	0.41

```
> findAssocs(dtm, "port", 0.4)
```

port	ssh	password
1.00	0.96	0.62

Create a Dictionary

- Create a dictionary specifying a subset of terms to work with
- Use the Dictionary function.

```
> ##Other functions related to selecting specific terms
>
> #Create a dictionary: a subset of terms
> d = Dictionary(c("root", "richard", "invalid", "session"))
> #Use dictionary to make a matrix with only the dictionary terms
> dtm_dictionary = TermDocumentMatrix(data2, list(dictionary = d))
> show(dtm_dictionary)
```

A term-document matrix (4 terms, 1000 documents)

Non-/sparse entries: 757/3243

Sparsity : 81%

Maximal term length: 7

Weighting : term frequency (tf)

Outlier Detection

How to Read Outlier Exhibits

- Each slide has a table containing outlier(s), and a scatterplot graph
- This example code produces one slide for every outlier
- If the occurrence frequency of a term happened to be perfectly correlated to the number of log entries for each month, points would line up diagonally on the regression line.
- Outliers are clearly visible, shown as points that are far removed from the regression line.
- Terms shown above the line are those that had more occurrences than expected, given the number of log entries for that month; terms below had less.

Outlier Detection Method

- Bonferoni outlier detection method was used to identify outliers.
- The Bonferroni method is less sensitive to false positives than outlier detection based on percentage increases, standard deviations (i.e., z-score), or other methods.
 - Identifying outliers based on percentage increase thresholds produces far more false positives for small denominator values, and too few for large denominators.
 - The Standard deviation method reduces false positives somewhat, but not enough.
 - The Bonferroni method generally tends to be more balanced across all value ranges.
- The Bonferroni outlier table at the top of each slide contains statistically significant outlier(s).
- The Bonferroni method detects outliers based on t-test of regression coefficients, at the 99.9% confidence level.
- An additional filter was added so that only those states that changed by at least plus or minus 10% can be considered outliers.
- Without outlier detection, thousands of slides would be required to show all the combinations of peril type, business line, and financial division, for each and every data attribute.

Term: invalid

Month	Term	Term Frequency	Total Month's Entries
Nov	invalid	0	170




Figure: Comparison of term frequency, invalid, with number of log entries

Term: port

Month	Term	Term Frequency	Total Month's Entries
Nov	port	43	170




Figure: Comparison of term frequency, port, with number of log entries

Term: password

Month	Term	Term Frequency	Total Month's Entries
Feb	password	33	54




Figure: Comparison of term frequency, password, with number of log entries

Term: session

Month	Term	Term Frequency	Total Month's Entries
May	session	83	611
Nov	session	96	170




Figure: Comparison of term frequency, session, with number of log entries

Most Significant Outlier Terms and Months

- Can write a loop to summarize the most significant outlier terms in a table
- Same terms as shown in preceding scatterplot graphs

	Month	Term	Term Frequency	Total Month's Entries
2	Nov	invalid	0	170
3	Nov	port	43	170
4	Feb	password	33	54
5	May	session	83	611
6	Nov	session	96	170